

Something about Submitter Segmentation

Anna Marbut

May 2019

Executive Summary

This is where the executive summary goes

Contents

Executive Summary	1
1 Introduction	3
2 Background	4
2.1 Classification algorithms	4
2.2 Google PageRank	4
3 Data and Descriptive Statistics	6
3.1 Data Source	6
3.2 User Accounts	6
3.3 Opportunities and Submissions	8
3.4 Labels	10
4 Methodology/Analysis	12
5 Results	13
6 Conclusion	14
Appendices	16
.1 Label Lists	16

Introduction

Submittable was founded in 2010 by three software developers who were creatives when they were off the clock—one novelist, one musician, and one filmmaker. All three had experience with the submission process for creative work, and saw a need for improvement. This process varied for every opportunity: one organization might want to be mailed a physical copy of the work, another might want to be mailed a thumb drive, another might receive submissions via e-mail, and yet another might use a digital filesharing service. Submittable’s founders had experience with the frustrating variety of requirements on the submitter-side, but they quickly realized how much worse it was for the organizations receiving the submissions. These organizations were using physical filing systems, spreadsheets, e-mail chains, and any other number of quasi-structured methods to track the submissions that they’d received and the various communications, review processes, and administrative tasks related to them. This left a lot of room for human error and increased frustration on both the submitter and the organization sides of the whole process.

With these issues in mind, the three founders set out to create a platform which could accept all of the most common filetypes (even very large files) and streamline the entire submission management pipeline. Almost 10 years later, Submittable has helped 15 thousand organizations collect and manage over 12 million submissions. While the bulk of their clients remain in the creative realm (literary journals, film contests, photography journals, auditions, etc.), they have branched out into almost every industry in which any type of submission needs to be made (research grants, job applications, charitable foundations, residency programs, etc.), and the number of these other clients is growing quickly. The company has grown from three founders to 70 employees, was picked up by the startup incubator Y-Combinator (known for DropBox, AirBnB, Weebly, and others), and has received over \$5 million in venture capital funding. In short, the founders’ idea seems to have been a good one.

Submittable’s business model has been one of “software-as-a-service”, or SAAS, in which their clients pay a subscription fee for use of the software platform. Different subscription levels are available based on which features each organization wants to use. Some of these include the number of active opportunities allowed (also referred to as “calls to submission” or “forms”), the option for multiple reviewers, gallery voting (such as for public photography contests), submission fee collection, and automated reporting. As the business has grown, the list of features offered to client organizations has expanded from those immediately related to the software functionality, to more abstract features such as hosting an opportunity on a custom url, white-labelling the software, or providing marketing services to promote a specific opportunity or organization. This last feature is the inspiration and focus for this project.

- Labelling issues
- Benefits of submitter segmentation for company and submitters

Background

2.1 Classification algorithms

— supervised vs. unsupervised

— ADA clustering?

— machine learning

— SVM and other supervised techniques?

twitter user profiling: <https://www.aaai.org/ocs/index.php/ICWSM/ICWSM11/paper/viewFile/2886>

— gradient boosted decision trees

classifier comparisons: <https://ieeexplore.ieee.org/abstract/document/5170933>

— Naive bayesian, Instance-based learner, classification and regression tree, Naive Bayesian tree

twitter user classification: <https://dl.acm.org/citation.cfm?id=1871993>

— stacked-SVM

2.2 Google PageRank

twitter community identification: http://snap.stanford.edu/networks2013/papers/netnips2013_submission

— used pagerank to identify central communities and community members in twitter network

Multilabel user classification using the community structure of online networks:

<https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0173347>

— inferring user interests based on sparsely annotated user profiles

— use pagerank to identify most important community memberships for each user, then apply to graph embeddings

Personalized PageRank Estimation and Search: <https://arxiv.org/abs/1507.05999>

— bidirectional personalized pagerank; meant to increase speed of search, but may be useful in bidirectional network (ie. submitters w/ common opportunities)

Identifying Key Users in Online Social Networks: <https://epub.uni-regensburg.de/25914/1/Identifying>

— importance of connectivity in marketing based on social networks

— concept of eigenvector centrality to acknowledge inequality of importance in network nodes

— uses weighted vector to apply bias for high-low activity users; **bidirectional activity**

weights; address undirected connectivity

— used dampening = 0.85; supported that higher d does not drastically change order of high ranking results, but takes much longer

Edge-Weighted Personalized PageRank: <http://wenleix.github.io/paper/edgeppr.pdf>

— node vs. edge weights; ie importance of submitter vs importance of relationship between submitters; in our context, node weight could be seen as topicality of submitter (how many labeled opportunities submitted to, keywords in user description), edge weight could be seen as topicality of opportunity that submitters have in common (# of relevant labels, keywords in opportunity description)

— focuses on providing quick(er) solution for applying edge weights to pagerank

— use reduced model of edge-weighted pagerank to decrease computation time

Topic-Sensitive PageRank: <http://www-cs-students.stanford.edu/~taherh/papers/topic-sensitive-pagerank.pdf>

— use pre-computed topic-specific pagerank vectors to rank search results; avoids highly connected nodes from rising to top when not topical

- pagerank matrix must be irreducible (highly connected, limited clustering?) and aperiodic (no sinks, solved by damping factor)
- topic bias added @ damping factor: $\alpha \vec{p}$ where $\alpha = P(\text{teleportation})$ and $\vec{p} = 1/N$ or $\vec{p} =$ biased rank vector
- power iteration looks like: $\vec{Rank} = (1 - \alpha)M * \vec{Rank} + \alpha \vec{p}$; so that biased rank is added at each iteration? how does it ever reach convergence? Wouldn't there always be difference of damping factor?
- reiterates that value of α does not have significant effect on ordering of results; this study used $\alpha = 0.25$

Data and Descriptive Statistics

3.1 Data Source

The data for this project comes from an Amazon Redshift Database that Submittable has set up to perform basic analytic queries about their clients, users, and product usage. Due to the scope of this research, the data used is centralized around Submittable's users, including information that they've explicitly provided on their Submittable profile, as well as information about the various submissions that they've made. In accordance with Submittable's privacy policy, no information that could be used to identify an individual user was included in the data for this project.

3.2 User Accounts

Submittable has 4.2 million users from all time, 770K of whom made a submission in 2018. As seen in Figure 3.1, the number of new user accounts has increased exponentially, with almost twice as many new accounts in 2018 as there were in 2017. The number of users making at least one submission is also increasing steadily each year, although on a linear scale with about 100K more active submitters each year since 2013 (Figure 3.2). On average Submittable users have made five submissions in their account lifetimes, with the number of submissions that an individual user has made ranging from one to 7.4K. Additionally, 62% of submittable users have only made one submission since they created their account.

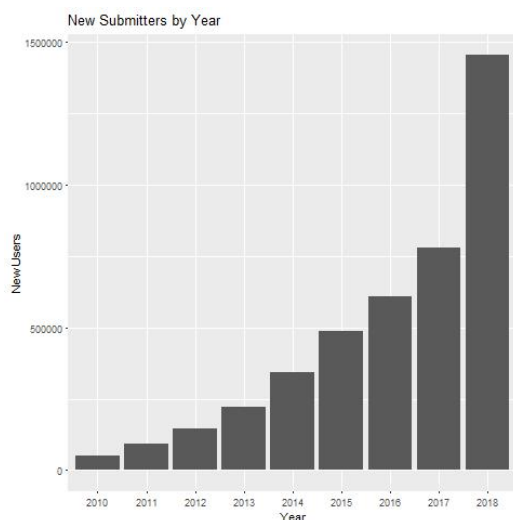


Figure 3.1: New Submittable User Accounts by Year




Figure 3.2: Submittable Users with at least One Submission by Year

In creating a Submittable account, each submitter has the option to provide a written description of themselves and various demographic information. Roughly 100K submitters have something written in their description field, and a brief scan shows that these are fairly descriptive, like what would be included in a cover letter or a professional bio. Only 850 of Submittable’s 4.2 million users have provided information about their gender on their Submittable profile, and only 760 have provided valid birthdates. Among these users, 72% self-reported as female, and the reported age distribution is as seen in Figure 3.3, with the overwhelming majority of users between 36 and 45 years old. 1.2 million of Submittable’s users have provided location information, hailing from 244 countries and territories. However, 75% of these users are from the United States, and 91% of them are from the top ten countries shown in Table3.1.

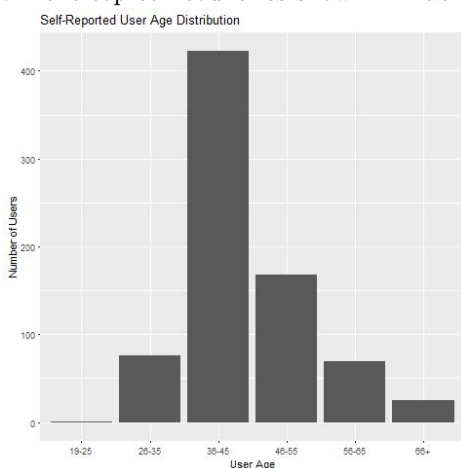


Figure 3.3: Submittable Users’ Ages

Table 3.1: Top Ten Submittable Users’ Countries

	Country	Number of Submitters
1	United States	913542
2	United Kingdom	59268
3	Canada	56316
4	Australia	26898
5	India	22709
6	Nigeria	7458
7	Germany	7359
8	France	5652
9	South Africa	5578
10	Italy	5451

3.3 Opportunities and Submissions

Submittable’s software platform has hosted a total of 130K opportunities, with a fairly steady increase of 3-5K new opportunities created every year except for 2018 (Figure 3.4). This decrease in new opportunities last year can likely be attributed to the maturation of many client accounts, who are now re-using forms that they’ve created in previous years. This is also supported by the steady increase in the total number of submissions made each year, including last year, which has increased by 300K every year since 2010 (Figure 3.5). In its lifetime, Submittable has received a total of 12M submissions.




Figure 3.4: New Submittable Opportunities by Year

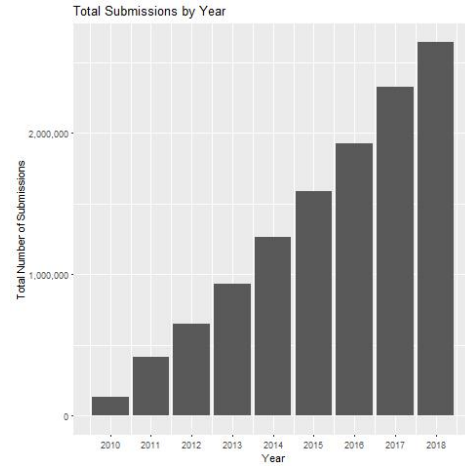


Figure 3.5: Total Number of Submissions Made by Year

On average, opportunities which have received at least one submission have received 125 submissions, with a range from one to 145K. About 75% of opportunities have received at least one submission, 45% have received at least five submissions, and 30% have received over 20 submissions. The intermittent nature of many opportunities makes it difficult to define what should be considered an “active” opportunity, but Figure 3.6 shows the number of opportunities that received at least one submission each year. These numbers follow a similar, though less dramatic, pattern as the number of new opportunities each year (Figure 3.4), steadily increasing until 2018.

Each opportunity has various textual data associated with it, including the opportunity name, a brief description of the opportunity, and any number of form field descriptions. A formfield description can be anything that the organization wants to put on their online form, from something as simple as “Name:” to a full writing prompt or detailed submission instructions. The opportunity descriptions also vary greatly in their descriptive usefulness, some consisting of multiple paragraphs describing the opportunity guidelines in detail, and others simply stating “Complete the form below”.




Figure 3.6: Number of Opportunities Receiving at least One Submission by Year

3.4 Labels

As Submittable has grown, it has seen several generations of labeling methodologies. Currently, the only actively-updated labels are “usecase” labels applied at the organizational level, and “discover” labels applied at the opportunity level. The usecase labels are assigned by account managers and are meant to describe what type of opportunities the organization will offer, such as “publishing”, “fellowships”, or “grants”. The discover labels are assigned by the organizations themselves as they opt in to Submittable’s opportunity search engine (called “Discover”), and can be anything from a list of almost 200 options that they think submitters might search for in relation to their opportunity. The full lists of usecase and discover labels can be found in Appendix .1

Only about 40K opportunities are associated with a usecase label (ie. are hosted by an organization with a usecase label), but this makes sense since the labels only started being applied in a consistent manner in the last year, and there were roughly 40K active opportunities in 2018. The ten usecase labels with the highest number of associated opportunities are shown in Table 3.2.

67K discover labels are associated with opportunities, however, since the relationship between the discover label and an opportunity is many-to-one, only 20K unique opportunities are associated with discover labels. Also, since many of the discover labels are overlapping (ie. “book” vs. “memoir” vs. “stories”), I needed to map these onto broader category labels. Submittable identified grant applicants, filmmakers, photographers, and writers (with the subcategories of poetry, fiction, and non-fiction writers), so I grouped the discover labels under these categories. Table 3.3 shows the number of unique forms with discover labels that fall into these categories. Since we know that only 20K unique opportunities are associated with discover labels, this table (which sums to about 40K) reiterates the many-to-one relationship between discover labels and opportunities.

Table 3.2: Top Ten Usecase Labels

	Usecase	Count
1	Publishing	10836
2	Grants	7406
3	Contest	4938
4	Award/ Nomination	4283
5	Exhibition	2626
6	Conference	2342
7	Festival or Event	1879
8	Other	1631
9	Fellowship	1613
10	Scholarships	1311

Table 3.3: Grouped Discover Labels

	Discover Group	Count
1	Fiction	9871
2	Film	5302
3	Grant	3157
4	Nonfiction	8311
5	Photography	5451
6	Poetry	8100

Methodology/Analysis

- compare results to keyword match lists
 - if time: compare response rate (click-through or even email open) for pagerank-developed list vs. keyword match
 - discussion of cost balance for bigger vs. more targeted list for marketing purposes

Results

Conclusion

Appendices

.1 Label Lists

	Usecase
1	Publishing
2	Other
3	Admissions
4	Scholarships
5	Job applications
6	Fellowship
7	Festival or Event
8	Peer Review
9	Artwork submissions
10	Audition
11	Grant applications
12	Contest or Campaign Entries
13	Video/Audio submissions
14	Award/ Nomination
15	Contest
16	Grants
17	Conference
18	Exhibition
19	Residency
20	Manuscript/Content submissions
21	Internal Use
22	Fellowship applications
23	Internships
24	Festival or Event submissions
25	Corporate Giving
26	Conference submissions
27	General applications
28	Audition submissions

Discover.Label	Discover.Label	Discover.Label	
1 interviews	51 science	101 magazine	
2 book	52 criticism	102 anthology	
3 review	53 hybrid	103 writing	
4 art	54 dance	104 flash	
5 fiction	55 student	105 feedback	
6 prose	56 exhibition	106 short-film	
7 literary	57 philosophy	107 screenwriting	
8 short-story	58 military	108 feminist	
9 video	59 grant	109 young-adult	
10 submishmash	60 native-american	110 novel	
11 nonfiction	61 children	111 agent	
12 memoir	62 column	112 online	Discover.Label
13 essay	63 health	113 magical-realism	151 music
14 stories	64 politics	114 translation	152 canada
15 travel	65 culture	115 adventure	153 gallery
16 visual-art	66 social-justice	116 award	154 mythology
17 creative-writing	67 volunteer	117 residency	155 emerging
18 scripts	68 animation	118 plays	156 event
19 contest	69 speculative	119 haiku	157 music-composition
20 theme	70 subscribe	120 community	158 parenting
21 publishing	71 print	121 business	159 minorities
22 manuscript	72 festival	122 funding	160 crafts
23 science-fiction	73 multilingual	123 editor	161 printmaking
24 drawing	74 comedy	124 digital	162 workshop
25 painting	75 humor	125 book-arts	163 internship
26 photography	76 conference	126 journalism	164 novella
27 nonprofit	77 retreat	127 audio	165 mentorship
28 chapbook	78 environmental	128 surrealist	166 history
29 lyrics	79 scholarships	129 documentary	167 ceramics
30 critique	80 public-art	130 film	168 museum
31 sculpture	81 cover-art	131 mystery	169 orchestra
32 classes	82 independent	132 technology	170 new-york
33 experimental	83 workspace	133 comics	171 fiber-arts
34 blog	84 art-fair	134 pitch	172 paper
35 for-sale	85 african-american	135 illustration	173 poster
36 monologue	86 paranormal	136 religion	174 thriller
37 design	87 juried	137 fellowship	175 360-Video
38 sports	88 season	138 performance	176 tequila
39 podcast	89 asian-american	139 prize	177 service-learning
40 academic	90 alcohol	140 romance	
41 television	91 fashion	141 horror	
42 sound	92 midwest	142 education	
43 collaborative	93 VR/360	143 lgbtqia	
44 international	94 female	144 fantasy	
45 graphic-novel	95 VR	145 food	
46 article	96 public-media	146 architecture	
47 short-play	97 civil-rights	147 spoken-word	
48 theater	98 mental-health	148 indigenous	
49 erotica	99 journal	149 readings	
50 job	100 poetry	150 media	

Bibliography